

CASE STUDY OF INSURANCE.CSV

In [1]:

```
1 !pip install scipy
```

Requirement already satisfied: scipy in c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages (1.6.3)

WARNING: You are using pip version 21.1.1; however, version 21.3.1 is available.

You should consider upgrading via the 'c:\users\ritesh\appdata\local\program s\python\python39\python.exe -m pip install --upgrade pip' command.

Requirement already satisfied: numpy<1.23.0,>=1.16.5 in c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages (from scipy) (1.20.3)

1. Import neccessry libraries

In [2]:

```
1
2 import numpy as np
3 import pandas as pd
4 import seaborn as sns
5 import matplotlib.pyplot as plt
6 import os
```

2. Read the data as a data frame

In [3]:

```
1 data= pd.read_csv(r'D:Reports/insurance.csv')
2 df=pd.DataFrame(data)
3 print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

[1338 rows x 7 columns]

Number 3

a. Shape of the data

In [4]:

```
1 df.shape
```

Out[4]:

```
(1338, 7)
```

b. Data type of each attribute

In [5]:

```
1 df.dtypes
```

Out[5]:

```
age          int64
sex          object
bmi          float64
children     int64
smoker       object
region       object
charges      float64
dtype: object
```

c. Checking the presence of missing values

In [6]:

```
1 for i in df.columns: print(data[i].isnull().sum())
2
3
```

```
0
0
0
0
0
0
0
0
```

other method for checking presence missing value

In [7]:

```
1 numeric = df.columns[df.dtypes != 'object']
2 print(numeric)
```

```
Index(['age', 'bmi', 'children', 'charges'], dtype='object')
```

In [8]:

```
1 df[numeric].isnull()
```

Out[8]:

	age	bmi	children	charges
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
...
1333	False	False	False	False
1334	False	False	False	False
1335	False	False	False	False
1336	False	False	False	False
1337	False	False	False	False

1338 rows × 4 columns

In [9]:

```
1 character = df.columns[df.dtypes=='object']
2 print(character)
```

Index(['sex', 'smoker', 'region'], dtype='object')

In [10]:

```
1 df[character].isnull()
```

Out[10]:

	sex	smoker	region
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
...
1333	False	False	False
1334	False	False	False
1335	False	False	False
1336	False	False	False
1337	False	False	False

1338 rows × 3 columns

In [11]:

```
1 df[character].isnull().sum()
```

Out[11]:

```
sex      0
smoker   0
region   0
dtype: int64
```

In [12]:

```
1 df[numeric].isnull().sum()
```

Out[12]:

```
age      0
bmi      0
children 0
charges  0
dtype: int64
```

d.calculate a 5-number summary

In [13]:

```
1 df.describe()
```

Out[13]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

In [14]:

```
1 #for Age
2 print(df['age'].describe()[['min', '25%', '50%', '75%', 'max']])
3
4 #for BMI
5 print(df['bmi'].describe()[['min', '25%', '50%', '75%', 'max']])
6
7 #for Children
8 print(df['children'].describe()[['min', '25%', '50%', '75%', 'max']])
9
10 #for Charges
11 print(df['charges'].describe()[['min', '25%', '50%', '75%', 'max']])
12
13
```

```
min      18.0
25%      27.0
50%      39.0
75%      51.0
max      64.0
Name: age, dtype: float64
min      15.96000
25%      26.29625
50%      30.40000
75%      34.69375
max      53.13000
Name: bmi, dtype: float64
min       0.0
25%       0.0
50%       1.0
75%       2.0
max       5.0
Name: children, dtype: float64
min      1121.873900
25%      4740.287150
50%      9382.033000
75%     16639.912515
max     63770.428010
Name: charges, dtype: float64
```

e. Distribution of 'bmi', 'age' and 'charges' columns.

In [15]:


```
1 sns.distplot(df['bmi'])  
2
```

c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[15]:

<AxesSubplot:xlabel='bmi', ylabel='Density'>




In [16]:

```
1 sns.distplot(df['age'])
```

```
c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[16]:

<AxesSubplot:xlabel='age', ylabel='Density'>




In [17]:

```
1 sns.distplot(df['charges'])
```

c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

Out[17]:

```
<AxesSubplot:xlabel='charges', ylabel='Density'>
```



In [18]:

```
1 # Dis_bmi=(
2 #     df.bmi
3 #     .value_counts()
4 #     .value_counts()
5 #     .sort_index()
6 # Dis_bmi.plot()
```

In [19]:

```
1 # Dis_age=(
2 #     df.age
3 #     .value_counts()
4 #     .value_counts()
5 #     .sort_index()
6 # Dis_age.plot()
```

In [20]:

```
1 # Dis_charges=(
2 #     df.charges
3 #     .value_counts()
4 #     .value_counts()
5 #     .sort_index()
6 # Dis_charges.plot()
```

f. Measure of skewness of 'bmi', 'age' and 'charges' columns

In [21]:

```
1 from scipy.stats import skew
```

In [22]:

```
1 print(skew(df.bmi, bias= False))
```

0.28404711059874477

In [23]:

```
1 print(skew(df.age, bias= False))
```

0.055672515652991854

In [24]:

```
1 print(skew(df.charges, bias= False))
```

1.5158796580240383

g. Checking the presence of outliers in 'bmi', 'age' and 'charges' columns

In [25]:

```
1 # import warnings
2 # warnings.filterwarnings('ignore')
3 # plt.figure(figsize=(16,5))
4 # plt.subplot(1,2,1)
5 # sns.distplot(df['bmi'])
6 # plt.subplot(1,2,2)
7 # sns.distplot(df['age'])
8 # plt.show()
```

In [26]:

```
1 # print("Highest allowed",df['bmi'].mean() + 3*df['bmi'].std())
2 # print("Lowest allowed",df['bmi'].mean() - 3*df['bmi'].std())
```

In [27]:


```
1 sns.boxplot(df['bmi'])
```

```
c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
warnings.warn(
```

Out[27]:

```
<AxesSubplot:xlabel='bmi'>
```



In [28]:


```
1 sns.boxplot(df['age'])
```

```
c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
warnings.warn(
```

Out[28]:

<AxesSubplot:xlabel='age'>



In [29]:


```
1 sns.boxplot(df['charges'])
```

c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[29]:

```
<AxesSubplot:xlabel='charges'>
```



h. Distribution of categorical columns (include children)

In [30]:

```
1 x=df.iloc[:,:].values
2 y=df.iloc[:,6].values
3 print(x)
4 print(y)
```


```
[[19 'female' 27.9 ... 'yes' 'southwest' 16884.924]
 [18 'male' 33.77 ... 'no' 'southeast' 1725.5523]
 [28 'male' 33.0 ... 'no' 'southeast' 4449.462]
 ...
 [18 'female' 36.85 ... 'no' 'southeast' 1629.8335]
 [21 'female' 25.8 ... 'no' 'southwest' 2007.945]
 [61 'female' 29.07 ... 'yes' 'northwest' 29141.3603]]
[16884.924 1725.5523 4449.462 ... 1629.8335 2007.945 29141.3603]
```

In [31]:

```
1 plt.bar(df['smoker'],y,)
2 plt.xlabel('smoker')
3 plt.ylabel('Charges')
4 plt.title('smoker vs charges')
5 plt.figure()
```

Out[31]:


<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>


In [32]:

```
1 plt.figure(figsize=(10,6))
2 sns.boxplot(x='children',y='age',data=df)
3 plt.show()
```



In [33]:


```
1 plt.figure(figsize=(10,6))
2 sns.countplot(x='smoker',data=df,palette='Blues')
3 plt.show()
```



In [34]:

```
1 plt.figure(figsize=(14,8))
2 sns.relplot(x='age',y='bmi',data=df)
3 plt.show()
```

<Figure size 1008x576 with 0 Axes>



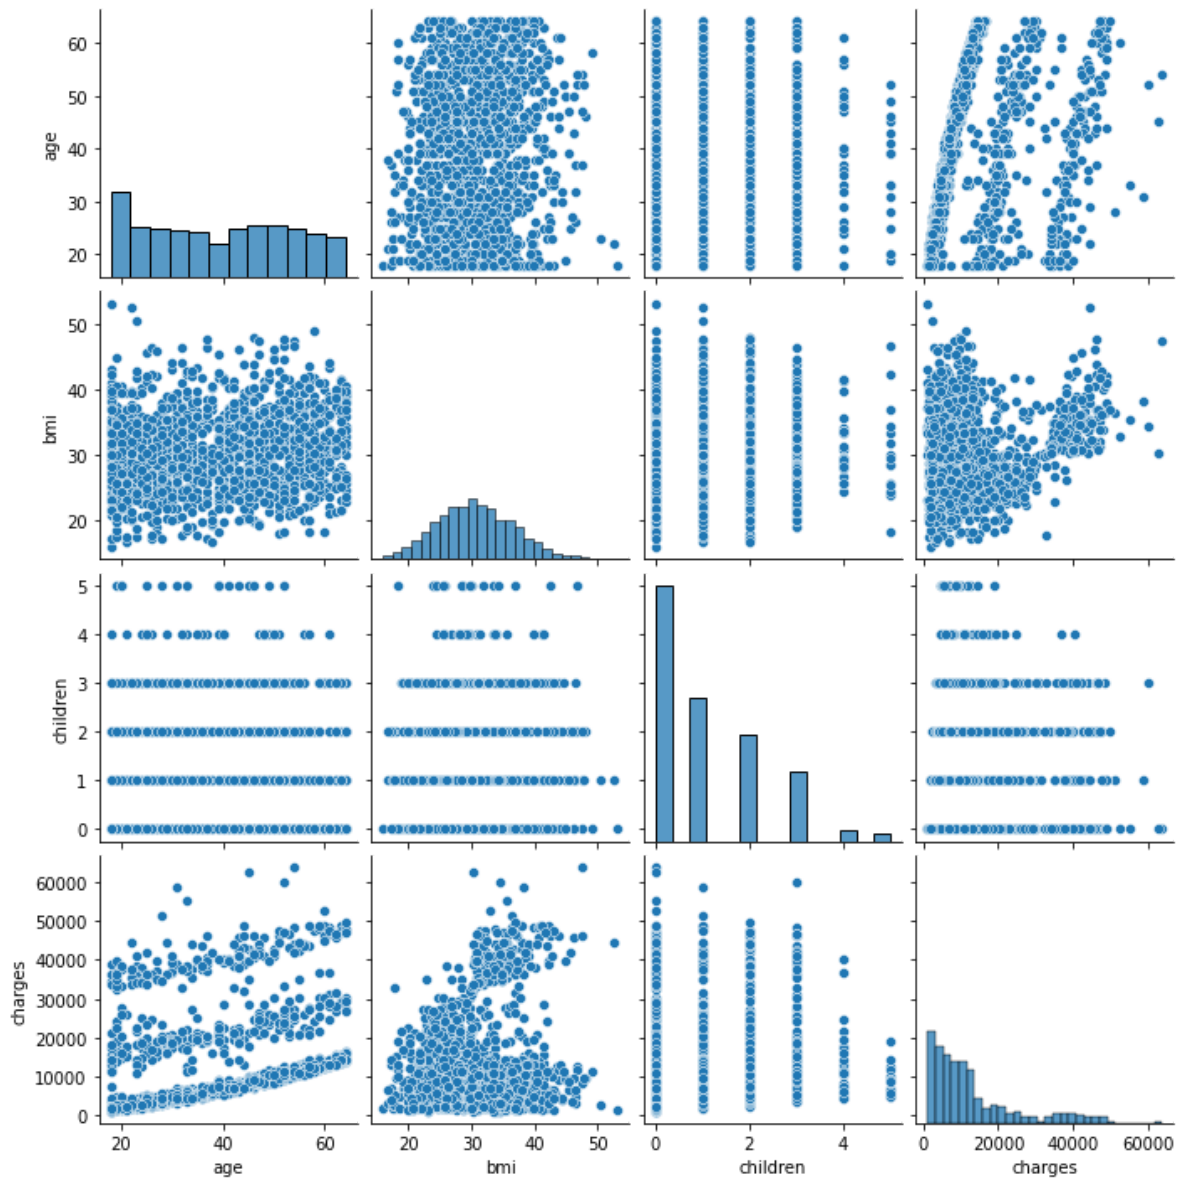
i. Pair plot that includes all the columns of the data frame

In [35]:

```
1 sns.pairplot(df)
```

Out[35]:

<seaborn.axisgrid.PairGrid at 0x1a65d044e80>

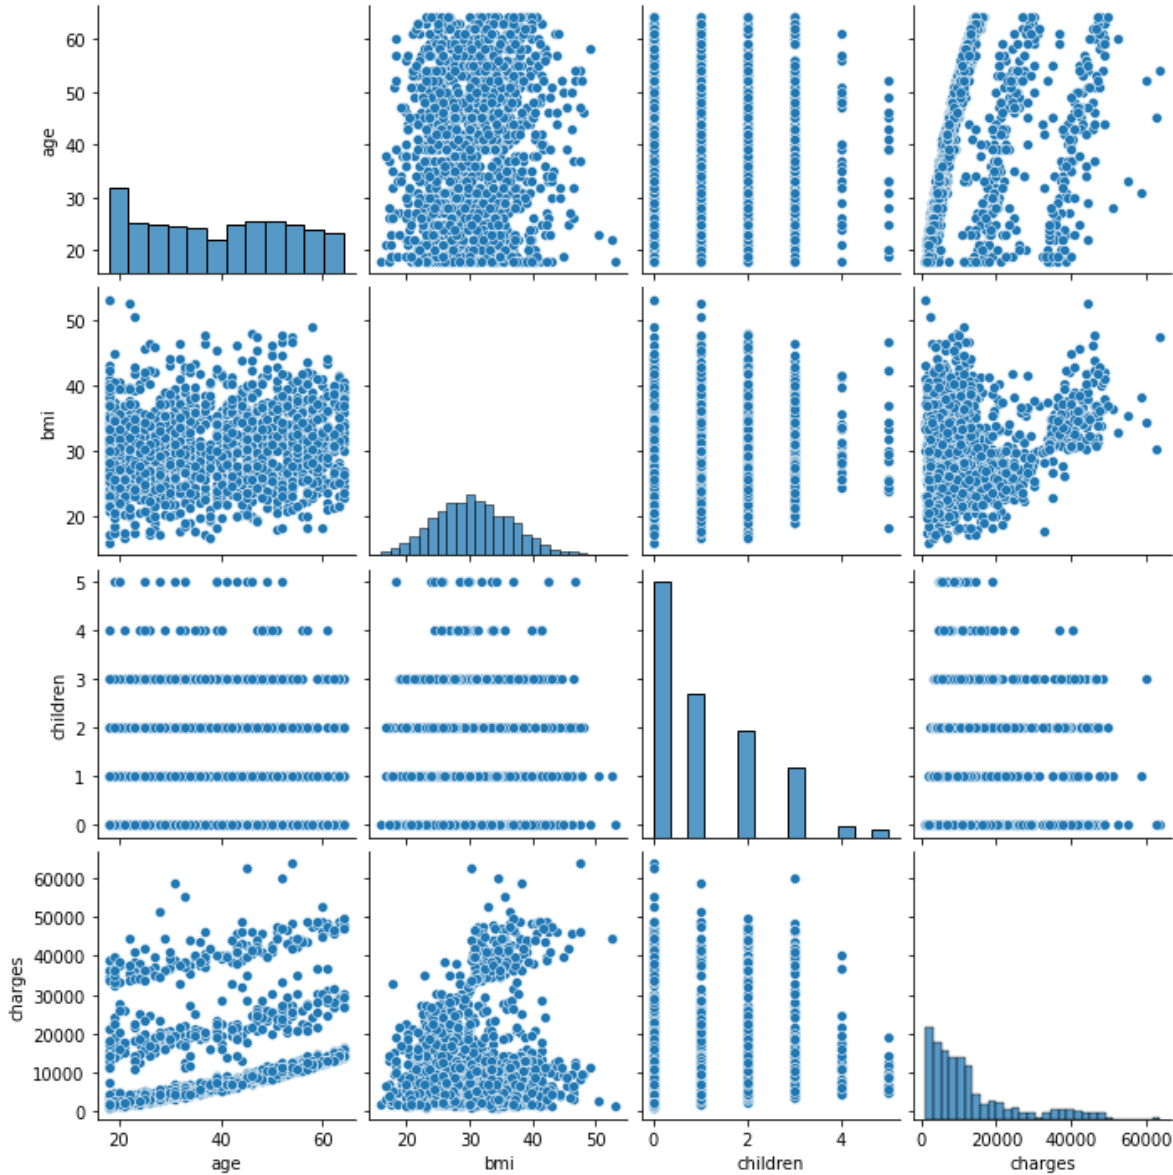


In [36]:

```
1 sns.pairplot(df[['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges']])
2
```

Out[36]:

<seaborn.axisgrid.PairGrid at 0x1a65d11d7c0>



As there are only 4 attributes numeric in nature , the rest of them are not included in the pair plot.

The diagonal plots represent distribution of the variable relative to itself . The plot between 'age' and 'bmi' does not indicate any relationship between them. The plot is more like a cloud and lacks direction. Between 'age' and 'charges', there is no clear relationship, though there seem to be 3 lines of positive relationship between them. It means, there are 3 sets of charges which increase gradually with age. No clear relation between 'age' and 'children' either. The range of 'bmi' decreases as children increases, however there are some extreme values in 'bmi' for children value 5.

Number 4

a. Do charges of people who smoke differ significantly from the people who don't?

People who smoke have high charges and the people who do not smoke have considerably very low charges. So, charges do differ for people who smoke from the people who do not smoke but not significantly as there is some intersection of values for both types of people.


In [37]:

```
1 sns.stripplot(data['charges'], data['smoker'])
2
```

```
c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

Out[37]:

```
<AxesSubplot:xlabel='charges', ylabel='smoker'>
```



In [41]:

```
1 df.smoker.value_counts()
```

Out[41]:

```
no      1064
yes      274
Name: smoker, dtype: int64
```

b. Does bmi of males differ significantly from that of females?

There is no significant difference in BMI for male and female genders, so no relationship exists between the two.

In [38]:


```
1 sns.stripplot(data['bmi'], data['sex'])
2
```

```
c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
warnings.warn(
```

Out[38]:

```
<AxesSubplot:xlabel='bmi', ylabel='sex'>
```



In [42]:

```
1 df.sex.value_counts()
```

Out[42]:

```
male      676
female    662
Name: sex, dtype: int64
```

c. Is the proportion of smokers significantly different in different genders?

The proportions being 58% and 42% for male and female genders who smoke are not significantly different.

In [39]:

```

Total number of smokers is --", df[df['smoker']=='yes'].shape[0])
Total number of male smokers is --", df[df['smoker']=='yes'][df['sex']=='male'].shape[0])
Total count of female smokers is -- ", df[df['smoker']=='yes'][df['sex']=='female'].shape[0])
Proportion of smokers who are male is --", (df[df['smoker']=='yes'][df['sex']=='male'].shape[0])/df[df['smoker']=='yes'].shape[0])
Proportion of smokers who are female is -- ", (df[df['smoker']=='yes'][df['sex']=='female'].shape[0])/df[df['smoker']=='yes'].shape[0])

```

```

Total number of smokers is -- 274
Total number of male smokers is -- 159
Total count of female smokers is -- 115
Proportion of smokers who are male is -- 0.5802919708029197
Proportion of smokers who are female is -- 0.4197080291970803

```

```

<ipython-input-39-9869b121467f>:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```

```

print("Total number of male smokers is --", df[df['smoker']=='yes'][df['sex']=='male'].shape[0])

```

```

<ipython-input-39-9869b121467f>:3: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```

```

print("Total count of female smokers is -- ", df[df['smoker']=='yes'][df['sex']=='female'].shape[0])

```

```

<ipython-input-39-9869b121467f>:4: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```

```

print("Proportion of smokers who are male is --", (df[df['smoker']=='yes'][df['sex']=='male'].shape[0])/df[df['smoker']=='yes'].shape[0])

```

```

<ipython-input-39-9869b121467f>:5: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```

```

print("Proportion of smokers who are female is -- ", (df[df['smoker']=='yes'][df['sex']=='female'].shape[0])/df[df['smoker']=='yes'].shape[0])

```

d. Is the distribution of bmi across women with no children, one child and two children, the same?

Yes, the distributions of 'bmi' are nearly same across women with 0, 1 or 2 children.

In [40]:


```
1 sns.stripplot(df['bmi'], df[df['sex']=='female']['children'])
```

c:\users\ritesh\appdata\local\programs\python\python39\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

Out[40]:

<AxesSubplot:xlabel='bmi', ylabel='children'>



Number 5

In [62]:

```

1 from sklearn.datasets import make_classification
2
3 X,Y=make_classification()
4 X,Y

```

Out[62]:

```

(array([[ -0.35346275,  0.11470425, -0.76527075, ..., -1.30705054,
         0.8945163 , -0.16308364],
       [ 0.45616974,  0.45554952,  0.69630583, ..., -0.25875887,
         1.83840011,  1.91229045],
       [ 0.60698229,  2.01889117,  0.60645728, ..., -0.0211031 ,
         0.10671858, -0.91777694],
       ...,
       [-0.11235949,  0.26517082, -1.025135 , ...,  0.64714497,
        -1.95875912,  0.68709827],
       [-0.44086909, -1.6436062 ,  0.31800784, ...,  2.69867185,
         0.89515976, -0.5175695 ],
       [-0.61268517, -1.80563453, -1.19828744, ..., -1.39129918,
        -1.32068727, -0.89897423]]),
array([1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0,
       0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0,
       1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1,
       0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1,
       0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0]))

```

In [66]:

```

1 print("The shape of X is {} and Y is {}".format(X.shape,Y.shape))

```

The shape of X is (1338, 6) and Y is (1338,)

In [67]:

```

1 np.unique(Y)

```

Out[67]:

```

array([ 1121.8739 , 1131.5066 , 1135.9407 , ..., 60021.39897,
        62592.87309, 63770.42801])

```

In [71]:

```

1 data = pd.DataFrame(np.c_[X,Y])

```


In [72]:

```
1 data
```

Out[72]:

	0	1	2	3	4	5	6
0	19.0	1.0	27.900	0.0	0.0	1.0	16884.92400
1	18.0	0.0	33.770	1.0	1.0	0.0	1725.55230
2	28.0	0.0	33.000	3.0	1.0	0.0	4449.46200
3	33.0	0.0	22.705	0.0	1.0	3.0	21984.47061
4	32.0	0.0	28.880	0.0	1.0	3.0	3866.85520
...
1333	50.0	0.0	30.970	3.0	1.0	3.0	10600.54830
1334	18.0	1.0	31.920	0.0	1.0	2.0	2205.98080
1335	18.0	1.0	36.850	0.0	1.0	0.0	1629.83350
1336	21.0	1.0	25.800	0.0	1.0	1.0	2007.94500
1337	61.0	1.0	29.070	0.0	0.0	3.0	29141.36030

1338 rows × 7 columns

In []:

```
1
```

Number 6

Data Pre-Processing

Encoding the categorical features

In [63]:

```
1 # encoding sex column
2 df.replace({'sex':{'male':0,'female':1}}, inplace=True)
3
4 3 # encoding 'smoker' column
5 df.replace({'smoker':{'yes':0,'no':1}}, inplace=True)
6
7 # encoding 'region' column
8 df.replace({'region':{'southeast':0,'southwest':1,'northeast':2,'northwest':3}}, inplace=True)
```

In [64]:

```
1 X = df.drop(columns='charges', axis=1)
2 Y = df['charges']
```

In [65]:

```
1 print(X)
```

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	1
1	18	0	33.770	1	1	0
2	28	0	33.000	3	1	0
3	33	0	22.705	0	1	3
4	32	0	28.880	0	1	3
...
1333	50	0	30.970	3	1	3
1334	18	1	31.920	0	1	2
1335	18	1	36.850	0	1	0
1336	21	1	25.800	0	1	1
1337	61	1	29.070	0	0	3

[1338 rows x 6 columns]

In [47]:

```
1 print(Y)
```

```
0    16884.92400
1    1725.55230
2    4449.46200
3    21984.47061
4    3866.85520
...
1333  10600.54830
1334   2205.98080
1335   1629.83350
1336   2007.94500
1337  29141.36030
```

Name: charges, Length: 1338, dtype: float64

In [48]:

```
1 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

In [49]:

```
1 print(X.shape, X_train.shape, X_test.shape)
```

(1338, 6) (1070, 6) (268, 6)

Model Training

Linear Regression

In [50]:

```
1 regressor = LinearRegression()
```

In [51]:

```
1 regressor.fit(X_train, Y_train)
```

Out[51]:

LinearRegression()

In [52]:

```
1 LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

Out[52]:

LinearRegression()

Model Evaluation

In [53]:

```
1 # prediction on training data
2 training_data_prediction = regressor.predict(X_train)
```

In [54]:

```
1 # R squared value
2 r2_train = metrics.r2_score(Y_train, training_data_prediction)
3 print('R squared vale : ', r2_train)
```

R squared vale : 0.751505643411174

In [55]:

```
1 # prediction on test data
2 test_data_prediction = regressor.predict(X_test)
```

In [56]:

```
1 r2_test = metrics.r2_score(Y_test, test_data_prediction)
2 print('R squared vale : ', r2_test)
```

R squared vale : 0.7447273869684077

In [59]:

```
1 input_data = (31,1,25.74,0,1,0)
2
3 # changing input_data to a numpy array
4 input_data_as_numpy_array = np.asarray(input_data)
5
6 # reshape the array
7 input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
8
9 prediction = regressor.predict(input_data_reshaped)
10 print(prediction)
11
12 print('The insurance cost is USD ', prediction[0])
```

[3760.0805765]

The insurance cost is USD 3760.0805764960514

In []:

1